

Parameterization, Implementation & Optimization

Parameterization

1. Centralized Configuration - PipelineConfig Class

All pipeline parameters are declared in a Python `dataclass`, enabling consistent and clean access throughout the application.

```
@dataclass
class PipelineConfig:
    aws_access_key: str = None
    aws_secret_key: str = None
    bucket_name: str = "default-bucket"
    s3_directory: str = "sensor_data/"
    local_path: str = "/tmp/data"
    db_credentials_path: str = "/etc/db_creds.json"
    input_files: List[str] = None
    sensor_patterns: List[str] = None
    default_start_date: str = "2024-01-01"
    lookback_days: int = 30
    jdbc_fetch_size: int = 10000
    jdbc_num_partitions: int = 8
    write_mode: str = "overwrite"
    use_aqe: bool = True
    enable_skew_handling: bool = True
```

2. Parameter Sources

Source	Use Case
Airflow Variables	Runtime overrides (e.g., file paths, S3 dirs)
AWS Secrets Manager	Credentials (secure)
Local JSON Config	Developer/test config

3. Parameter Usage

- **DataLoader:** `bucket_name, s3_directory, input_files, local_path`
- **DataProcessor:** `sensor_patterns`
- **DatabaseManager:** `jdbc_fetch_size, jdbc_num_partitions`
- **S3Writer:** `write_mode`
- **Spark Session:** `use_aqe, enable_skew_handling`

4. Parameter Hierarchy

1. **Secrets Manager** → Highest priority
2. **Airflow Variables** → Mid-level overrides
3. **Local JSON** → Default fallback

5. Validation Checks

Before execution:

```
assert self.config.jdbc_fetch_size > 0
assert self.config.write_mode in ["overwrite", "append"]
assert all(f.endswith('.parquet') for f in self.config.input_files)
```

Section 2: Implementation Guide –

1. Infrastructure Setup

- **EC2 Instance:**
 - Dev: `t3.xlarge`
 - Prod: `r5.2xlarge`

- **IAM Role:** Attach S3 + SecretsManager + Glue permissions

2. Software Installation

Java + Python

```
sudo yum install java-11-amazon-corretto python3-pip git awscli -y  
pip3 install --upgrade pip
```

Spark

```
wget https://d1cdn.apache.org/spark/spark-3.4.1/spark-3.4.1-bin-hadoop3.tgz  
tar -xvf spark-3.4.1-bin-hadoop3.tgz  
sudo mv spark-3.4.1-bin-hadoop3 /opt/spark  
echo 'export SPARK_HOME=/opt/spark' >> ~/.bashrc  
echo 'export PATH=$PATH:$SPARK_HOME/bin' >> ~/.bashrc  
source ~/.bashrc
```

Python packages

```
pip install pyspark==3.4.1 boto3 pyarrow findspark
```

3. Configuration Files

/etc/db_creds.json

```
{  
  "YourDB": {  
    "host": "your-rds-endpoint",  
    "dbname": "sensor_db",  
    "user": "spark_user",  
    "password": "your_password"  
  }  
}
```

4. Deploy and Run

A. Clone & Navigate

```
git clone https://github.com/your-org/sensor-pipeline.git  
cd sensor-pipeline
```

B. Execute Locally

```
spark-submit --master local[4] main.py --config-source file --config-path  
config/pipeline_config.json
```

C. Execute on EC2

```
nohup spark-submit \  
  --master local[*] \  
  --executor-memory 4G \  
  --driver-memory 2G \  
  main.py --config-source aws --secret-name prod/sensor_pipeline \  
  > logs/run.log 2>&1 &
```

6. Monitoring & Logs

- **Spark UI:** http://<EC2_PUBLIC_IP>:4040

Log tail:

```
tail -f logs/run.log
```

Enable Spark event logs:

```
--conf spark.eventLog.enabled=true \  
--conf spark.eventLog.dir=s3://your-bucket/spark-logs/
```

7. Security Considerations

- Use IAM roles over keys wherever possible
 - Restrict security group to known IPs
 - Encrypt S3 + Secrets Manager
 - Rotate secrets regularly
-

Recommended and Applied Optimizations

1. Adaptive Query Execution (AQE)

Parameter: `use_aqe = True`

Spark Config: `.config("spark.sql.adaptive.enabled", self.config.use_aqe)`

- Dynamically adjusts the number of shuffle partitions.
 - Converts sort-merge joins to broadcast joins where applicable.
 - Significantly improves performance on large and skewed datasets.
-

2. Skew Join Handling

Parameter: `enable_skew_handling = True`

Spark Config: `.config("spark.sql.adaptive.skewJoin.enabled", self.config.enable_skew_handling)`

- Detects and mitigates data skew during shuffle-intensive joins.
 - Helps avoid long-running or failed stages due to uneven partition sizes.
-

3. Broadcast Join for Small Tables

Code: `df.join(broadcast(tags_df), df.tagid == tags_df.id, "left")`

- Broadcasts the smaller `tags_df` to all worker nodes.
 - Prevents shuffle joins, reducing network I/O and execution time.
-

4. Repartitioning for File Size Optimization

Code: `df.repartition(max(1, df.count() // 100000))`

- Dynamically repartitions DataFrame before writing to S3.

- Ensures approximately 100,000 records per file, balancing performance and read efficiency.
-

5. Use of `.persist(StorageLevel.MEMORY_AND_DISK)` for Caching

Code: `data[filename] = df.persist(StorageLevel.MEMORY_AND_DISK)`

- Caches frequently accessed DataFrames in memory with disk fallback.
 - Reduces redundant reads from S3 or local storage.
-

6. Data Coalescing During Reads

Code: `if df.rdd.getNumPartitions() > 1:`

`df = df.coalesce(1)`

- Reduces the number of partitions for small input files.
 - Avoids unnecessary parallelism that could degrade performance.
-

7. JDBC Parallelism

Parameters:

- `jdbc_fetch_size = 10000`
- `jdbc_num_partitions = 8`

Code:

```
"fetchSize": str(self.config.jdbc_fetch_size),  
"numPartitions": str(self.config.jdbc_num_partitions),  
"partitionColumn": "tagid"
```

- Enables **parallel reading** from PostgreSQL using the `tagid` column.

- Improves throughput and reduces bottlenecks during data ingestion.
-

8. File Size Capping via `maxRecordsPerFile`

Code: `.option("maxRecordsPerFile", 100000)`

- Prevents oversized Parquet files.
 - Optimizes read performance on S3 and improves downstream processing.
-

9. Duplicate Prevention with `dropDuplicates()`

Code: `sensor_df.dropDuplicates(["datetime"])`

- Ensures idempotency by removing duplicate entries before writing.
 - Prevents redundancy in historical sensor records.
-

10. Dynamic Output Partitioning and Unioning

Logic:

- Appends to existing data when `write_mode = "append"`.
 - Performs union with existing S3 files to maintain continuity.
-

11. Minimal Selective Projection and Column Casting

Code: `df.select([field.name for field in expected_schema])`

- Applies schema enforcement and column pruning.
 - Improves memory efficiency by avoiding wide transformations.
-

12. Clean Resource Unpersisting & Cache Management

Code: `df.unpersist()`

`self.spark.catalog.clearCache()`

- Ensures memory is released post-pipeline execution.
 - Prevents memory leaks during long-running or repeated jobs.
-